

Revenue growth divisions.

TYU division

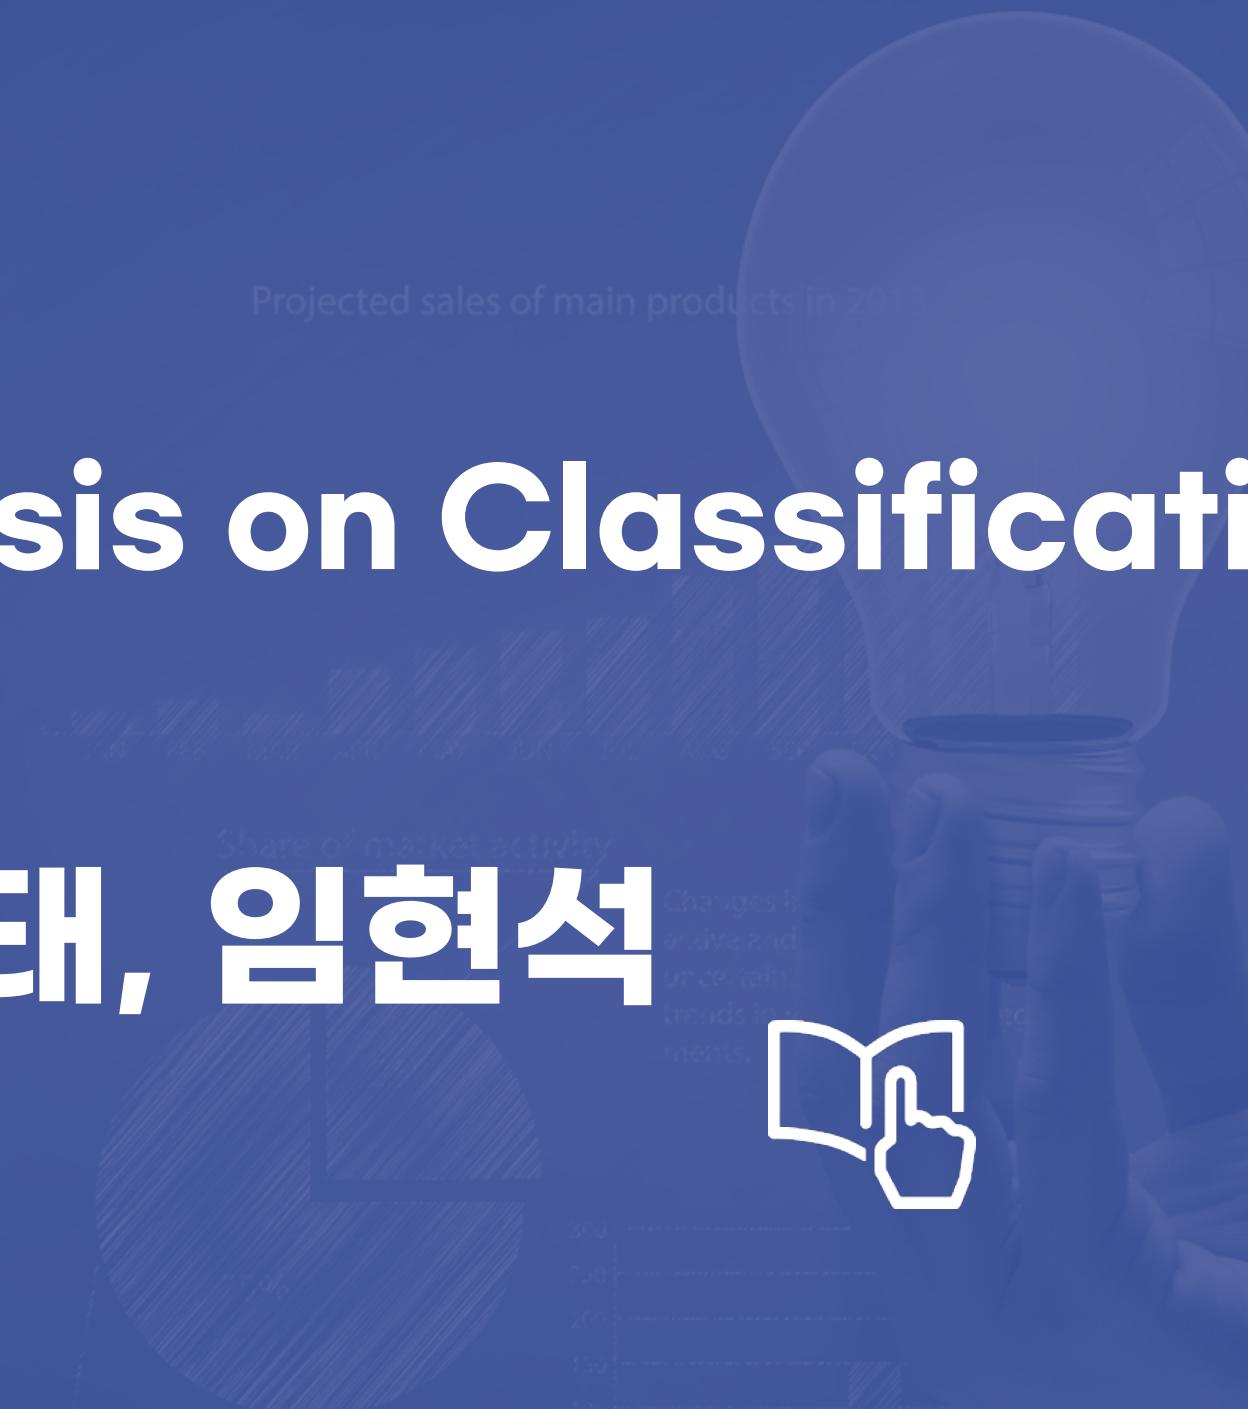
FRT division

Projected sales of main products in 2013

Sentiment Analysis on Classification

MLP LAB 임경태, 임현석

	TYU division			FRT division		
GHT	254	550	254	274	154	415
RDW	650	320	754	273	825	154
TRG	241	450	144	364	954	174
RTG	254	650	874	657	125	274
WCF	794	145	124	752	741	759
HBT	452	794	954	341	741	345



Contents

01 Logistic Regression

02 Logistic Regression codes

03 Sentiment Analysis

04 Logistic Classification

05 Logistic Classification codes

06 Sentiment Analysis of Restaurant Review

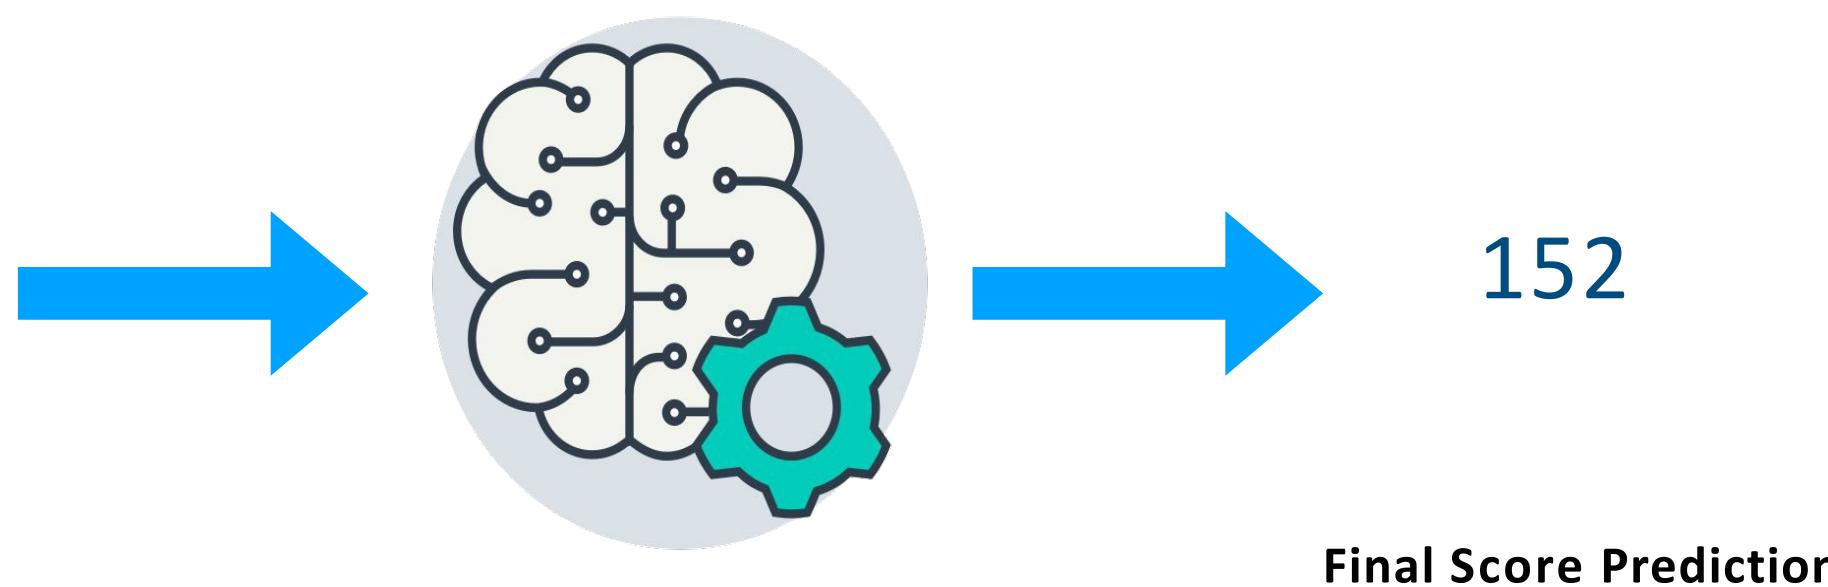


01. Logistic Regression

Recap

- Multi-variable Linear Regression

Quiz1: 73
Quiz2: 80
Quiz3: 75
Quiz Scores



$$H(x) = Wx + b$$

Questions

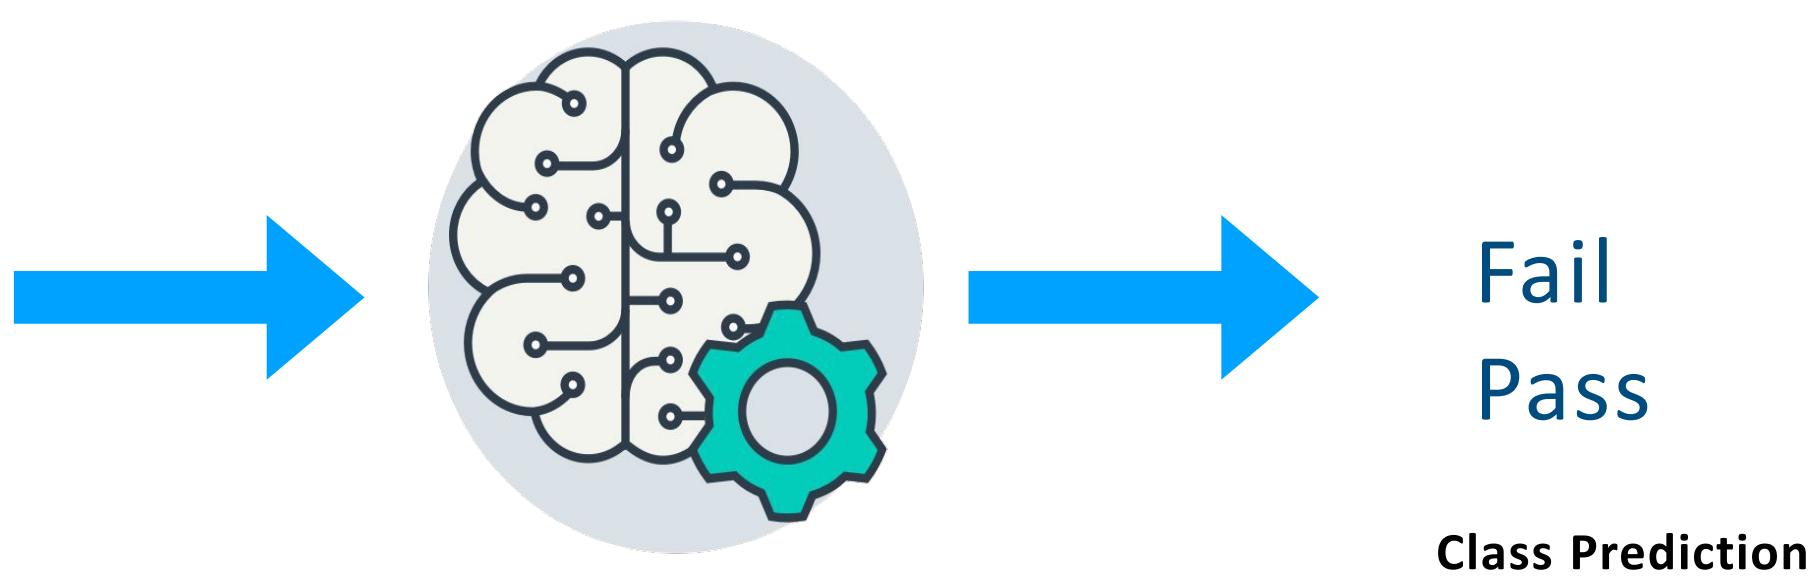
- How do I solve these cases?

Quiz1: 73

Quiz2: 80

Quiz3: 75

Quiz Scores



$$H(x) = Wx + b$$

Solution

- Logistic Regression (LR) Model
 - Let's switch from regression to a classification problem!

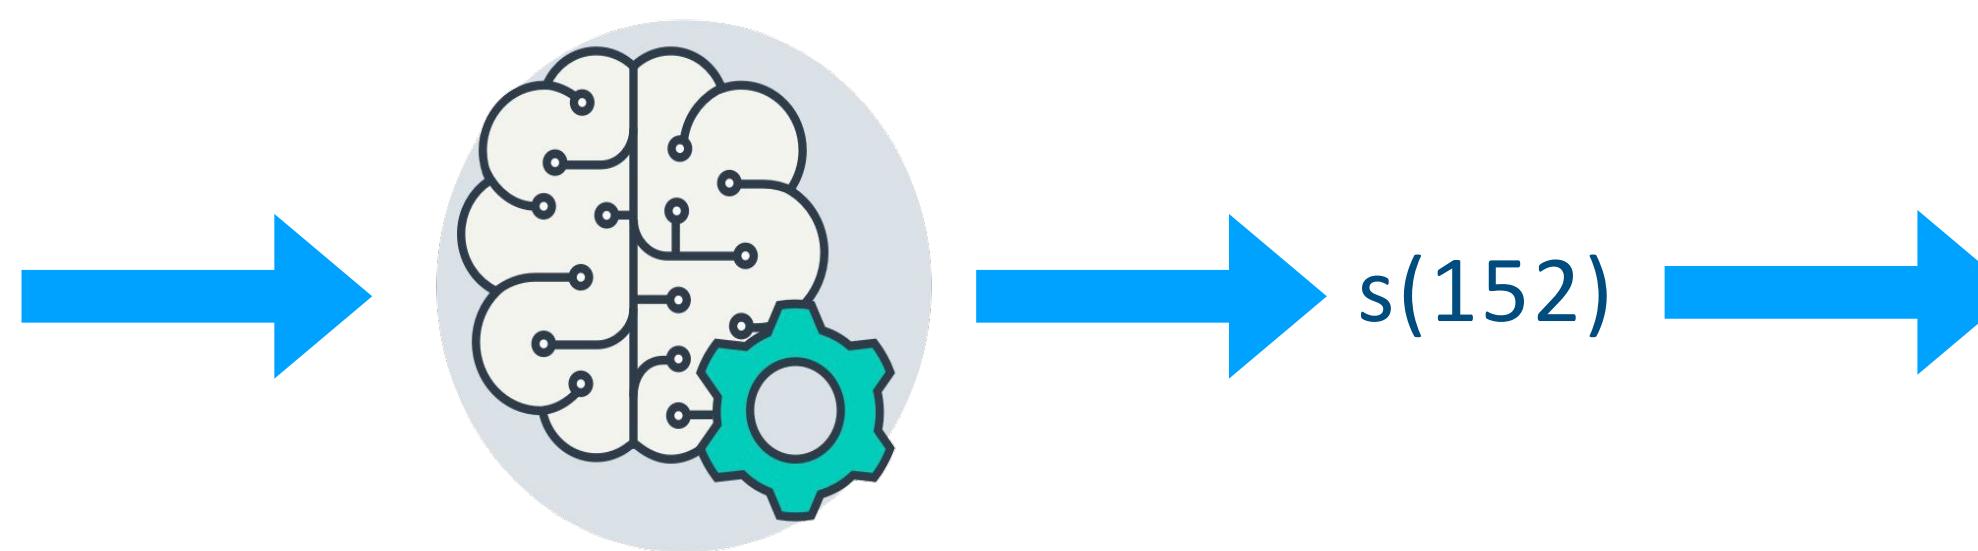
Simply output the model's output as 1 for passing and 0 for failing!

Quiz1: 73

Quiz2: 80

Quiz3: 75

Quiz Scores



$$H(x) = Wx + b$$

$$H(x) = s(Wx + b)$$

Logistic Regression (LR)

- Sigmoid function
 - A Function to convert from regression to classification problems

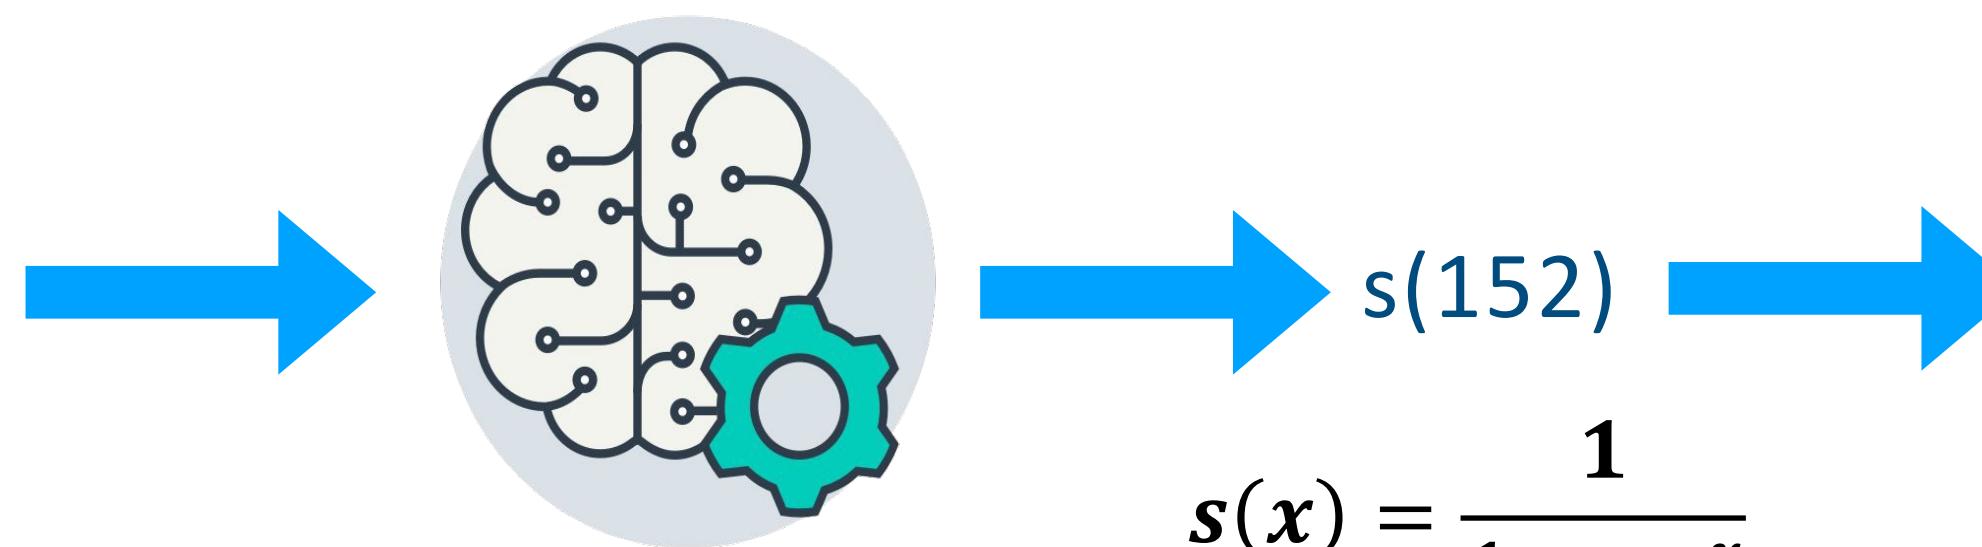
Change the output of a model to a value between 0 and 1

Quiz1: 73

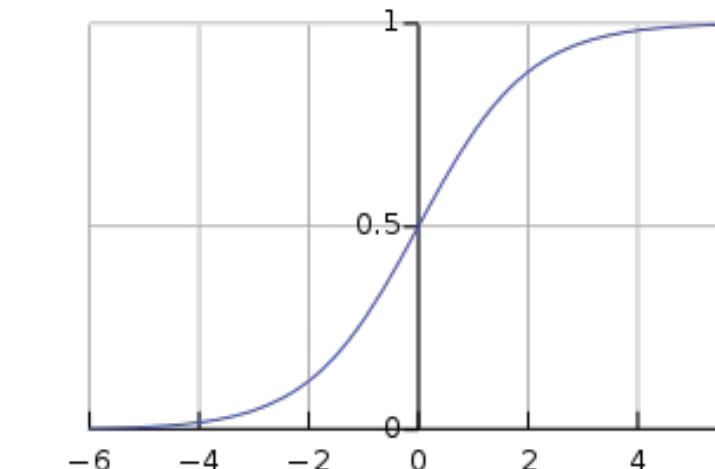
Quiz2: 80

Quiz3: 75

Quiz Scores



$$s(x) = \frac{1}{1 + e^{-x}}$$



$$H(x) = \frac{1}{1 + e^{-(Wx+b)}}$$

Logistic Regression (LR)

- Loss of LR?: Negative Log Likely-hood cost function

$$\text{- } Cost(W, b) = -\frac{1}{m} \sum_{i=1}^m y \log(H_\theta(x)) + (1 - y) \log(1 - H_\theta(x))$$

$$\left. \begin{aligned} \mathbb{P}(y = 1|z) &= \sigma(z) = \frac{1}{1 + e^{-z}} \\ \mathbb{P}(y = 0|z) &= 1 - \sigma(z) = \frac{1}{1 + e^z} \end{aligned} \right\} \mathbb{P}(y|z) = \sigma(z)^y (1 - \sigma(z))^{1-y}$$

Logistic Regression (LR)

- Gradient Descent Function

Weight Update via Gradient Descent

$$W := W - \alpha \frac{\partial}{\partial W} cost(W)$$

-
- α : Learning rate

02. Logistic Regression codes

Logistic Regression (LR) with Pytorch

- Data preparation

```
x_data = [[1, 2], [2, 3], [3, 1], [4, 3], [5, 3], [6, 2]]  
y_data = [[0], [0], [0], [1], [1], [1]]
```

Consider the following classification problem: given the number of hours each student spent watching the lecture and working in the code lab, predict whether the student passed or failed a course. For example, the first (index 0) student watched the lecture for 1 hour and spent 2 hours in the lab session ([1, 2]), and ended up failing the course ([0]).

```
x_train = torch.FloatTensor(x_data)  
y_train = torch.FloatTensor(y_data)
```

As always, we need these data to be in `torch.Tensor` format, so we convert them.

```
print(x_train.shape)  
print(y_train.shape)
```

```
torch.Size([6, 2])  
torch.Size([6, 1])
```

Logistic Regression (LR) with Pytorch

- Computing the Hypothesis

$$H(X) = \frac{1}{1 + e^{-W^T X}}$$

PyTorch has a `torch.exp()` function that resembles the exponential function.

```
print('e^1 equals: ', torch.exp(torch.FloatTensor([1])))  
e^1 equals:  tensor([2.7183])
```

We can use it to compute the hypothesis function conveniently.

```
w = torch.zeros((2, 1), requires_grad=True)  
b = torch.zeros(1, requires_grad=True)  
  
hypothesis = 1 / (1 + torch.exp(-(x_train.matmul(w) + b)))  
  
print(hypothesis)  
print(hypothesis.shape)  
  
tensor([[0.5000],  
        [0.5000],  
        [0.5000],  
        [0.5000],  
        [0.5000],  
        [0.5000]], grad_fn=<MulBackward0>)  
torch.Size([6, 1])
```

Logistic Regression (LR) with Pytorch

- Computing the Cost Function

$$cost(W) = -\frac{1}{m} \sum y \log(H(x)) + (1 - y) \log(1 - H(x))$$

We want to measure the difference between `hypothesis` and `y_train`.

```
print(hypothesis)
print(y_train)

tensor([[0.5000],
        [0.5000],
        [0.5000],
        [0.5000],
        [0.5000],
        [0.5000]], grad_fn=<SigmoidBackward>)
tensor([[0.],
        [0.],
        [0.],
        [1.],
        [1.],
        [1.]])
```

Logistic Regression (LR) with Pytorch

- Computing the Cost Function

For one element, the loss can be computed as follows:

```
- (y_train[0] * torch.log(hypothesis[0]) +  
  (1 - y_train[0]) * torch.log(1 - hypothesis[0]))  
  
tensor([0.6931], grad_fn=<NegBackward>)
```

Logistic Regression (LR) with Pytorch

- Computing the Cost Function from Scratch

To compute the losses for the entire batch, we can simply input the entire vector.

```
losses = -(y_train * torch.log(hypothesis) +
           (1 - y_train) * torch.log(1 - hypothesis))
print(losses)

tensor([0.6931],
      [0.6931],
      [0.6931],
      [0.6931],
      [0.6931],
      [0.6931], grad_fn=<NegBackward>)
```

Then, we just `.mean()` to take the mean of these individual losses.

```
cost = losses.mean()
print(cost)

tensor(0.6931, grad_fn=<MeanBackward1>)
```

Logistic Regression (LR) with Pytorch

- Computing the Cost Function using the Pytorch Functional package

```
F.binary_cross_entropy(hypothesis, y_train)  
tensor(0.6931, grad_fn=<BinaryCrossEntropyBackward>)
```

Logistic Regression (LR) with Pytorch

- Whole Training Procedure (w/o Class)

```
# 모델 초기화
W = torch.zeros((2, 1), requires_grad=True)
b = torch.zeros(1, requires_grad=True)
# optimizer 설정
optimizer = optim.SGD([W, b], lr=1)

nb_epochs = 1000
for epoch in range(nb_epochs + 1):

    # Cost 계산
    hypothesis = torch.sigmoid(x_train.matmul(W) + b) # or .mm or @
    cost = F.binary_cross_entropy(hypothesis, y_train)

    # cost로 H(x) 개선
    optimizer.zero_grad()
    cost.backward()
    optimizer.step()

    # 100번마다 로그 출력
    if epoch % 100 == 0:
        print('Epoch {:4d}/{}, Cost: {:.6f}'.format(
            epoch, nb_epochs, cost.item()
        ))
```

```
Epoch    0/1000 Cost: 0.693147
Epoch   100/1000 Cost: 0.134722
Epoch   200/1000 Cost: 0.080643
Epoch   300/1000 Cost: 0.057900
Epoch   400/1000 Cost: 0.045300
Epoch   500/1000 Cost: 0.037261
Epoch   600/1000 Cost: 0.031672
Epoch   700/1000 Cost: 0.027556
Epoch   800/1000 Cost: 0.024394
Epoch   900/1000 Cost: 0.021888
Epoch  1000/1000 Cost: 0.019852
```

Logistic Regression (LR) with Pytorch

- Evaluation

After we finish training the model, we want to check how well our model fits the training set.

```
hypothesis = torch.sigmoid(x_train.matmul(w) + b)
print(hypothesis[:5])

tensor([[0.4103],
       [0.9242],
       [0.2300],
       [0.9411],
       [0.1772]], grad_fn=<SliceBackward>)
```

Logistic Regression (LR) with Pytorch

- Evaluation

We can change **hypothesis** (real number from 0 to 1) to **binary predictions** (either 0 or 1) by comparing them to 0.5.

```
prediction = hypothesis >= torch.FloatTensor([0.5])
print(prediction[:5])

tensor([[0],
       [1],
       [0],
       [1],
       [0]], dtype=torch.uint8)
```

Logistic Regression (LR) with Pytorch

- Evaluation

Then, we compare it with the correct labels `y_train`.

```
print(prediction[:5])
print(y_train[:5])

tensor([[0],
       [1],
       [0],
       [1],
       [0]], dtype=torch.uint8)
tensor([[0.],
       [1.],
       [0.],
       [1.],
       [0.]])
```

Logistic Regression (LR) with Pytorch

- Evaluation

```
correct_prediction = prediction.float() == y_train
print(correct_prediction[:5])

tensor([[1],
       [1],
       [1],
       [1],
       [1]], dtype=torch.uint8)
```

Logistic Regression (LR) with Pytorch

- Higher Implementation with Class

```
class BinaryClassifier(nn.Module):
    def __init__(self):
        super().__init__()
        self.linear = nn.Linear(8, 1)
        self.sigmoid = nn.Sigmoid()

    def forward(self, x):
        return self.sigmoid(self.linear(x))

model = BinaryClassifier()
```

Logistic Regression (LR) with Pytorch

- Higher Implementation with Class

```
# optimizer 설정
optimizer = optim.SGD(model.parameters(), lr=1)

nb_epochs = 100
for epoch in range(nb_epochs + 1):

    # H(x) 계산
    hypothesis = model(x_train)

    # cost 계산
    cost = F.binary_cross_entropy(hypothesis, y_train)

    # cost로 H(x) 개선
    optimizer.zero_grad()
    cost.backward()
    optimizer.step()

    # 20번마다 로그 출력
    if epoch % 10 == 0:
        prediction = hypothesis >= torch.FloatTensor([0.5])
        correct_prediction = prediction.float() == y_train
        accuracy = correct_prediction.sum().item() / len(correct_prediction)
        print('Epoch {:4d}/{} Cost: {:.6f} Accuracy {:.2f}%'.format(
            epoch, nb_epochs, cost.item(), accuracy * 100,
        ))
```

```
Epoch  0/100 Cost: 0.704829 Accuracy 45.72%
Epoch 10/100 Cost: 0.572391 Accuracy 67.59%
Epoch 20/100 Cost: 0.539563 Accuracy 73.25%
Epoch 30/100 Cost: 0.520042 Accuracy 75.89%
Epoch 40/100 Cost: 0.507561 Accuracy 76.15%
Epoch 50/100 Cost: 0.499125 Accuracy 76.42%
Epoch 60/100 Cost: 0.493177 Accuracy 77.21%
Epoch 70/100 Cost: 0.488846 Accuracy 76.81%
Epoch 80/100 Cost: 0.485612 Accuracy 76.28%
Epoch 90/100 Cost: 0.483146 Accuracy 76.55%
```

03. Sentiment Analysis

Sentiment Analysis

- Sentiment analysis: A technique for categorizing the polarity of sentences and documents, such as positive/negative/neutral.

 Apple 아이패드 에어 5세대 M1 WIFI 64G 스페이스 그레이 (MM9C3KH…
최저 837,000원 무료배송 최저가 사려가기

판매처 94 제품정보 블로그리뷰 **쇼핑몰리뷰 1,445** AiITEMS추천

★★★★★ 5 하이마트쇼핑몰 · r*u***** · 22.07.06.
오래전부터 갖고 싶었던 아이패드
오래전부터 갖고 싶었던 아이패드.
몇일간 검색해본 결과 하이마트쇼핑몰에서 사는게 가격도 저렴하고 믿을 수가 있다는 생각에 구입 결정!
제품의 품질은 역시 애플 제품답게 훌륭하네요.
앞으로 제가 잘 쓰기만 하면 될 듯.
무엇보다 포장에 대단히 신경을 써주신 것 같아 좋았습니다. 제품이 박스에서 충격받지 않도록 종이 포장으로 칭칭 단단하게 감겨 있어 포장에 감동받기



리뷰펼치기 ▼

★★★★★ 5 하이마트쇼핑몰 · u*u***** · 23.01.29.
딸아이가 공부.인강들을때 필요
딸아이가 공부.인강들을때 필요하셔서 구입. 배송은 이틀만에 왔구요. 안포장은 잘되있는데...겉포장박스가 허술해서...조금그랬구요...패드는 하자없는
제품으로 잘 도착했어요. 카드청구할인으로 조금 저렴하게 구입했어요. 추가로 필요한 애플케어랑...액서사리값이 비싸네요. 그레이색상 질리지않고 무난
하니 잘 선택했구요.필름추가구입해서 바로 붙여서 사용했어요.

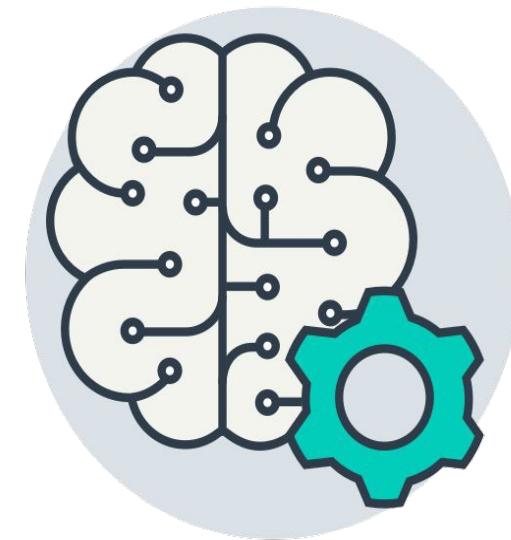


A simple Sentiment Analysis on LR

- Simple Modeling of a Binary Classification Problem

아이폰은 디자인이 예뻐
(iPhone is beautifully designed)

대화, 혹은 댓글



$s(152)$



Negative == 0
Positive == 1

Class Prediction

A simple Sentiment Analysis on LR

- Simple Modeling of a Binary Classification Problem
 - An example of “Bag of Words Meets Bags of Popcorn” at Kaggle
 - Can we refactor AWE Linear Regression code?

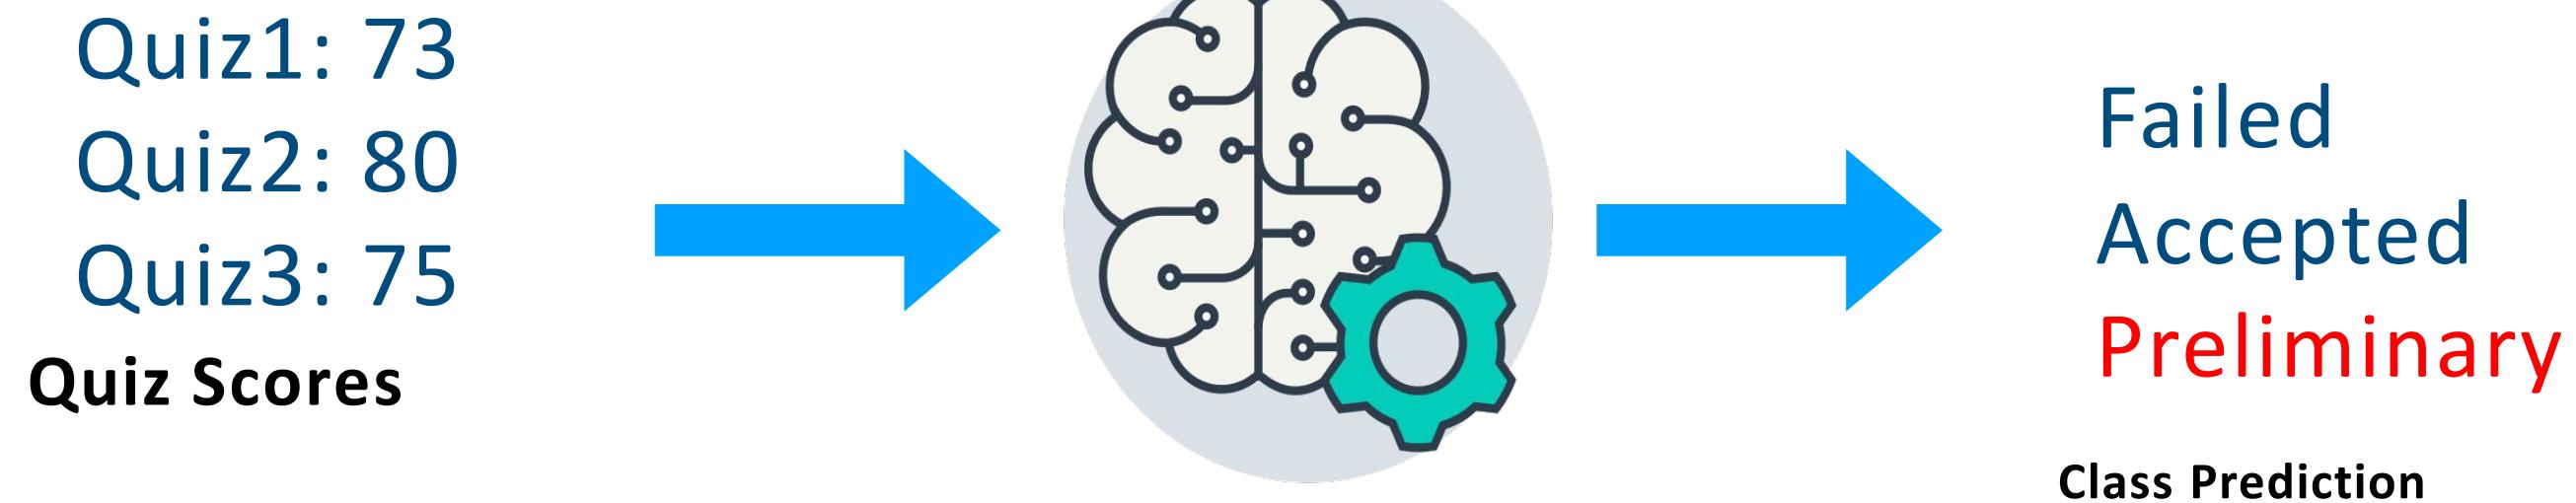
<https://www.kaggle.com/competitions/word2vec-nlp-tutorial/overview>

	id	sentiment	review
0	5814_8	1	With all this stuff going down at the moment w...
1	2381_9	1	\The Classic War of the Worlds\" by Timothy Hi...
2	7759_3	0	The film starts with a manager (Nicholas Bell)...
3	3630_4	0	It must be assumed that those who praised this...
4	9495_8	1	Superbly trashy and wondrously unpretentious 8...

04. Multinomial Classification

Questions

- How do I solve multi-class prediction?
 - This is known as the Multinomial Classification problem.



$$H(x) = Wx + b$$

Solution

- Probability Distribution Output with Softmax

- Let's switch from regression to a classification problem!

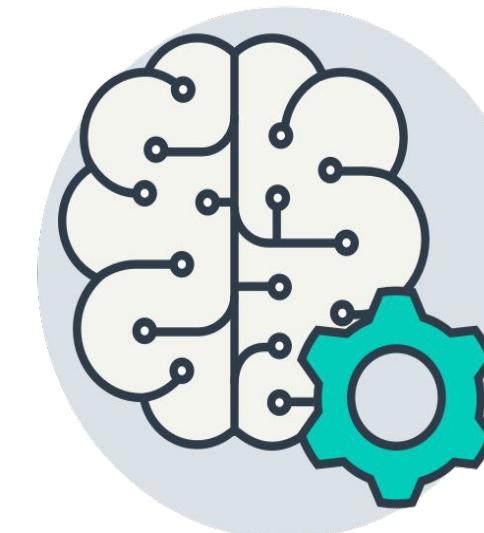
Simply make the output of the model a probability distribution over the number of classes.

Quiz1: 73

Quiz2: 80

Quiz3: 75

Quiz Scores



$s(11, 25, 37.1)$

$$\text{softmax}(\mathbf{x})_i = \frac{e^{x_i}}{\sum_j e^{x_j}}$$

Failed == 0.1

Accepted == 0.2

Preliminary == 0.7

Class Prediction



$$H(\mathbf{x}) = s(W\mathbf{x} + \mathbf{b})$$

Multinomial Classification

- Loss of Multinomial Classification: Cross Entropy
 - A method to calculate the distance between two probability distributions

$$H(P, Q) = -\mathbb{E}_{x \sim P(x)}[\log Q(x)] = -\sum_{x \in \mathcal{X}} P(x) \log Q(x)$$

$$- CE(\mathbf{y}, \hat{\mathbf{y}}) = - \sum_{i=1}^{N_c} y_i \log(\hat{y}_i)$$

```
import torch
def cross_ent_loss(x,y):
    return (y * -torch.log(x))

a = torch.FloatTensor([[0.1, 0.2, 0.7]])
b = torch.FloatTensor([[0, 0, 1]])
cross_ent_loss(a,b)
```

▼ 1. Basic usage of the CE loss

```
[62] import torch
     import matplotlib.pyplot as plt
     import numpy as np

[63] def cross_ent_loss(x,y):
     return (y * -torch.log(x))

[64] # if we have two different prob distributions below:
     a = torch.FloatTensor([[0.1, 0.2, 0.7]])
     b = torch.FloatTensor([[0, 0, 1]])
     cross_ent_loss(a,b)

[65] tensor([[0.0000, 0.0000, 0.3567]])

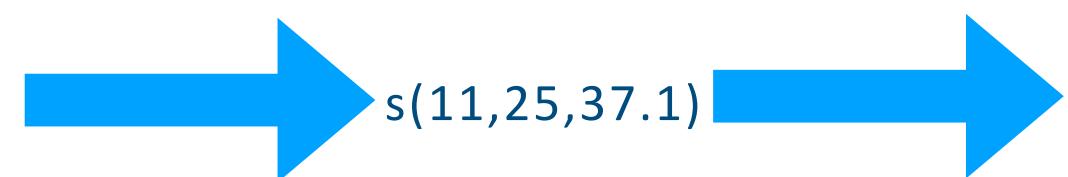
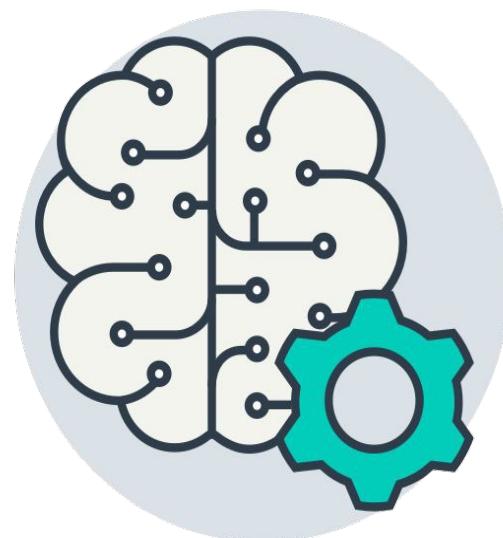
[66] # Even two distributions are almost identical, it returns marginal loss values
     a = torch.FloatTensor([[0, 0, 0, 0, 0.01, 0.99]])
     b = torch.FloatTensor([[0, 0, 0, 0, 0, 1]])
     cross_ent_loss(a,b)

[67] tensor([[    nan,    nan,    nan,    nan, 0.0000, 0.0101]])
```

Multinomial Classification

- Loss of Multinomial Classification: Cross Entropy
 - A method to calculate the distance between two probability distributions

$$CE(\mathbf{y}, \hat{\mathbf{y}}) = - \sum_{i=1}^{N_c} y_i \log(\hat{y}_i)$$



$$\text{softmax}(\mathbf{x})_i = \frac{e^{x_i}}{\sum_j e^{x_j}}$$

Failed == 0.1
Accepted == 0.2
Preliminary == 0.7
Class Prediction

CE loss

Failed == 0
Accepted == 1
Preliminary == 0
Target (실제정답)

Multinomial Classification

- Implementing Cross Entropy

Preparing the Data

- Let's create random dummy data and save it as z, assuming it results from a regression.
- To change the regression result to Multinomial Classification, we need Softmax!

For multi-class classification, we use the cross entropy loss.

$$L = \frac{1}{N} \sum -y \log(\hat{y})$$

where \hat{y} is the predicted probability and y is the correct probability (0 or 1).

```
z = torch.rand(3, 5, requires_grad=True)
hypothesis = F.softmax(z, dim=1)
print(hypothesis)

tensor([[0.2645, 0.1639, 0.1855, 0.2585, 0.1277],
        [0.2430, 0.1624, 0.2322, 0.1930, 0.1694],
        [0.2226, 0.1986, 0.2326, 0.1594, 0.1868]], grad_fn=<SoftmaxBackward0>)

y = torch.randint(5, (3,)).long()
print(y)

tensor([0, 2, 1])
```

Multinomial Classification

- Implementing Cross Entropy
 - Change the actual answer data to a probability distribution, one-hot encoding
- Calculate Loss: Calculate the CE Loss for the two probability distributions of the true answer and the predicted value

```
y_one_hot = torch.zeros_like(hypothesis)
y_one_hot.scatter_(1, y.unsqueeze(1), 1)

tensor([[1., 0., 0., 0., 0.],
        [0., 0., 1., 0., 0.],
        [0., 1., 0., 0., 0.]])  
  
cost = (y_one_hot * -torch.log(hypothesis)).sum(dim=1).mean()
print(cost)  
  
tensor(1.4689, grad_fn=<MeanBackward1>)
```

Multinomial Classification

- Cross Entropy (low, high-level)

```
# Low level
(y_one_hot * -torch.log(F.softmax(z, dim=1))).sum(dim=1).mean()

tensor(1.4689, grad_fn=<MeanBackward1>)
```

```
# High level
F.nll_loss(F.log_softmax(z, dim=1), y)

tensor(1.4689, grad_fn=<NllLossBackward>)
```

PyTorch also has `F.cross_entropy` that combines `F.log_softmax()` and `F.nll_loss()`.

```
F.cross_entropy(z, y)

tensor(1.4689, grad_fn=<NllLossBackward>)
```

Multinomial Classification

- Gradient Descent Function

Weight Update via Gradient Descent

$$W := W - \alpha \frac{\partial}{\partial W} cost(W)$$

-
- α : Learning rate

Thank you